

Complete Guide to I/O System Calls in C

File Descriptors, Standard Streams, and Low-Level I/O Operations

Understanding Linux/Unix System Programming

Systems Programming Masterclass

Compiled: January 15, 2026

Abstract

Abstract: This comprehensive guide explains Input/Output System Calls in C programming on Linux/Unix systems. We explore the fundamental concepts of file descriptors, standard streams, and the five essential I/O system calls: create, open, close, read, and write. Understanding these low-level operations is crucial for system programming, understanding how operating systems work, and writing efficient, low-level code.

Contents

1	Introduction to System Calls	4
1.1	Why Learn I/O System Calls?	4
1.2	Difference Between Library Functions and System Calls	4
1.3	Document Structure	4

2 File Descriptors and Standard Streams	5
2.1 What is a File Descriptor?	5
2.1.1 Key Characteristics	5
2.2 Kernel Data Structures for File I/O	5
2.2.1 Standard Streams in Action	7
2.3 File Descriptor Table Visualization	7
2.4 Key Concepts Summary	9
3 The 5 Essential I/O System Calls	10
3.1 1. Create System Call	10
3.1.1 Function Details	10
3.1.2 Permission Bits (mode)	10
3.1.3 <code>create()</code> Example	10
3.1.4 How <code>creat()</code> Works Internally	11
3.2 2. Open System Call	12
3.2.1 Function Details	12
3.2.2 Common Open Flags	12
3.2.3 <code>open()</code> Examples	12
3.2.4 Path Types	13
3.3 3. Close System Call	14
3.3.1 Function Details	14
3.3.2 <code>close()</code> Examples	14
3.3.3 What Happens on <code>close()</code> ?	15
3.3.4 Important <code>close()</code> Behavior	15
3.4 4. Read System Call	17
3.4.1 Function Details	17
3.4.2 <code>read()</code> Examples	17
3.4.3 Important <code>read()</code> Behaviors	18
3.4.4 <code>read()</code> Return Value Summary	19
3.4.5 Common <code>read()</code> Errors	19
3.5 5. Write System Call	20
3.5.1 Function Details	20
3.5.2 <code>write()</code> Examples	20
3.5.3 Important <code>write()</code> Behaviors	21
3.5.4 <code>write()</code> Return Value Summary	22
3.5.5 Common <code>write()</code> Errors	22
4 Practical Examples and Use Cases	24
4.1 Example 1: Simple File Copy	24
4.2 Example 2: Simple Cat Command	25
4.3 Example 3: Simple Echo Command	26
4.4 Example 4: Simple Head Command	27
5 Advanced Topics and Best Practices	29
5.1 Error Handling with <code>errno</code>	29
5.2 File Positioning with <code>lseek</code>	30
5.2.1 <code>lseek()</code> Examples	30
5.3 File Status with <code>fstat</code>	31
5.4 Duplicating File Descriptors with <code>dup/dup2</code>	33
5.5 Non-blocking I/O	34

6 Common Pitfalls and Debugging	35
6.1 Common System Call Errors	35
6.2 Debugging Techniques	35
6.3 Best Practices Summary	37
6.4 Performance Considerations	37
Conclusion	37

1 Introduction to System Calls

What Are System Calls?

System calls are interfaces between user programs and the operating system kernel. They allow programs to request services from the kernel for operations they don't have direct access to, such as:

- File operations (create, read, write, delete)
- Process management
- Memory allocation
- Device communication
- Network operations

1.1 Why Learn I/O System Calls?

- **Direct access** to operating system functionality
- **Complete control** over file operations
- **Understanding** how high-level I/O functions work internally
- **Essential for** system programming and embedded systems
- **Better performance** in certain scenarios
- **Fundamental knowledge** for Linux/Unix programming

1.2 Difference Between Library Functions and System Calls

Aspect	Library Functions	System Calls
Location	User-space libraries	Kernel-space
Performance	Faster (no context switch)	Slower (context switch required)
Examples	<code>printf()</code> , <code>fopen()</code>	<code>write()</code> , <code>open()</code>
Buffering	Usually buffered	Usually unbuffered
Flexibility	Higher level, easier to use	Lower level, more control
Portability	More portable	Less portable (OS-specific)

1.3 Document Structure

- **Section 2:** File Descriptors and Standard Streams
- **Section 3:** The 5 Essential I/O System Calls
- **Section 4:** Practical Examples and Use Cases
- **Section 5:** Advanced Topics and Best Practices
- **Section 6:** Common Pitfalls and Debugging

2 File Descriptors and Standard Streams

2.1 What is a File Descriptor?

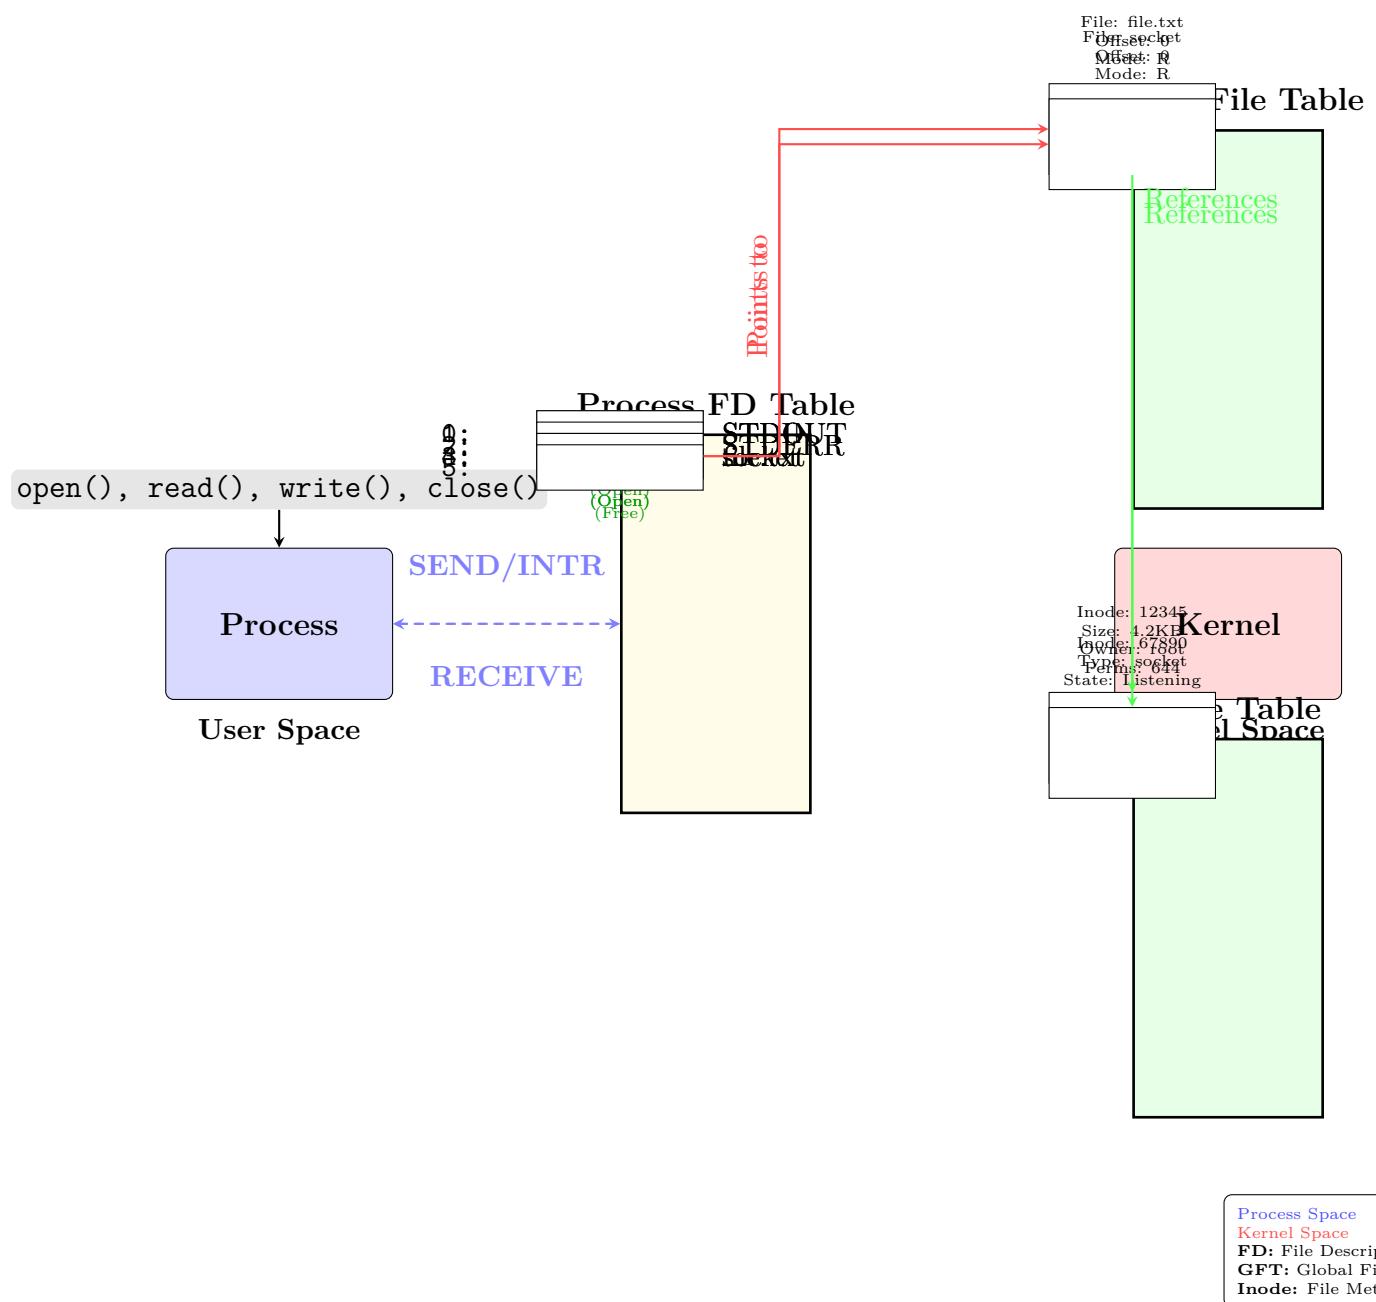
File Descriptor Definition

A **file descriptor** (FD) is a non-negative integer that uniquely identifies an opened file within a process. It's an index into the kernel's **file descriptor table** for that process.

2.1.1 Key Characteristics

- **Integer value** (0, 1, 2, 3, 4, ...)
- **Process-specific** (each process has its own FD table)
- **Kernel-managed** (maintained by the OS)
- **Lowest available** integer is always used for new files
- **Standard FDs** (0, 1, 2) are reserved for special purposes

2.2 Kernel Data Structures for File I/O



How File Descriptors Work

1. **System Call:** Process calls `open()` to request file access
2. **FD Allocation:** Kernel finds lowest free FD in Process FD Table
3. **GFT Entry:** Kernel creates entry in Global File Table with file state
4. **Inode Lookup:** GFT entry references file metadata in Inode Table
5. **Pointer Chain:** FD → GFT entry → Inode entry
6. **I/O Operations:** All `read()`/`write()` use FD to traverse chain
7. **Cleanup:** `close()` removes FD entry, may free GFT entry

Key Points:

- Each process has its own FD table
- Multiple FDs can point to same GFT entry (file sharing)
- GFT entries track per-open-file state (offset, mode)
- Inode table contains persistent file metadata

2.2.1 Standard Streams in Action

```

1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main() {
5     char buffer[100];
6
7     // Write to stdout (FD 1)
8     write(STDOUT_FILENO, "Enter your name: ", 17);
9
10    // Read from stdin (FD 0)
11    int bytes_read = read(STDIN_FILENO, buffer, sizeof(buffer)-1);
12
13    if (bytes_read > 0) {
14        buffer[bytes_read] = '\0';
15        // Write to stdout
16        write(STDOUT_FILENO, "You entered: ", 13);
17        write(STDOUT_FILENO, buffer, bytes_read);
18
19        // Write to stderr (FD 2)
20        write(STDERR_FILENO, "[DEBUG] Name captured\n", 22);
21    }
22
23    return 0;
24 }
```

Listing 1: Demonstrating Standard File Descriptors

2.3 File Descriptor Table Visualization

```

1 #include <stdio.h>
2 #include <unistd.h>
```

```

3 #include <fcntl.h>
4 #include <stdlib.h>
5
6 int main() {
7     int ifd, ofd, pipefd[2];
8
9     printf("Kernel Process File Descriptor Table\n");
10    printf("=====\\n");
11    printf("FD\\tType\\t\\tDescription\\n");
12    printf("-----\\n");
13
14    // 1. Standard FDs (0, 1, 2)
15    printf("%d\\tStandard\\tSTDIN (Keyboard/Input)\\n", STDIN_FILENO);
16    printf("%d\\tStandard\\tSTDOUT (Terminal/Output)\\n", STDOUT_FILENO);
17    printf("%d\\tStandard\\tSTDERR (Terminal/Error)\\n", STDERR_FILENO);
18
19    // 2. Open Files (Allocates 3 and 4)
20    ifd = open("in.txt", O_RDONLY | O_CREAT, 0644);
21    ofd = open("out.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644);
22
23    printf("%d\\tFile\\t\\tin.txt (Input File)\\n", ifd);
24    printf("%d\\tFile\\t\\tout.txt (Output File)\\n", ofd);
25
26    // 3. Create a Pipe (Allocates 5 and 6)
27    if (pipe(pipefd) == -1) {
28        perror("Pipe failed");
29        return 1;
30    }
31    printf("%d\\tPipe Read\\tInternal Buffer Read End\\n", pipefd[0]);
32    printf("%d\\tPipe Write\\tInternal Buffer Write End\\n", pipefd[1]);
33
34    // Cleanup
35    close(ifd); close(ofd);
36    close(pipefd[0]); close(pipefd[1]);
37
38    return 0;
39}

```

Listing 2: Visualizing File Descriptor Table

Expected Output

Kernel Process File Descriptor Table		
=====		
FD	Type	Description
0	Standard	STDIN (Keyboard/Input)
1	Standard	STDOUT (Terminal/Output)
2	Standard	STDERR (Terminal/Error)
3	File	in.txt (Input File)
4	File	out.txt (Output File)
5	Pipe Read	Internal Buffer Read End
6	Pipe Write	Internal Buffer Write End

2.4 Key Concepts Summary

Important Rules

- **Rule 1:** `open()` always returns the smallest unused file descriptor
- **Rule 2:** Standard FDs (0,1,2) are automatically opened for every process
- **Rule 3:** After `close()`, the FD becomes available for reuse
- **Rule 4:** File descriptors are inherited by child processes (`fork`)
- **Rule 5:** `dup()` and `dup2()` create copies of file descriptors

3 The 5 Essential I/O System Calls

3.1 1. Create System Call

create() - Create a New File

```
int creat(const char *pathname, mode_t mode);
```

3.1.1 Function Details

- **Header:** <fcntl.h> and <unistd.h>
- **Purpose:** Create a new empty file or truncate an existing file
- **Permissions:** mode specifies file permissions (octal)
- **Equivalent to:** open(pathname, O_WRONLY | O_CREAT | O_TRUNC, mode)
- **Limitation:** Always opens file for writing only

3.1.2 Permission Bits (mode)

Octal	Symbolic	Description
0400	r----	Owner can read
0200	-w---	Owner can write
0100	-x--	Owner can execute
0040	--r--	Group can read
0020	--w--	Group can write
0010	--x--	Group can execute
0004	---r-	Others can read
0002	---w-	Others can write
0001	---x-	Others can execute

3.1.3 create() Example

```

1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5
6 int main() {
7     // Create a new file with permissions rw-r--r--
8     int fd = creat("newfile.txt", 0644);
9
10    if (fd == -1) {
11        perror("Error creating file");
12        return 1;
13    }
14
15    printf("File 'newfile.txt' created successfully\n");
16    printf("File descriptor: %d\n", fd);
17

```

```
18 // Write some data to the file
19 char *data = "Hello, World!\n";
20 write(fd, data, 14);
21
22 // Close the file descriptor
23 close(fd);
24
25 return 0;
26 }
```

Listing 3: Using creat() System Call

3.1.4 How creat() Works Internally

1. Kernel checks if file already exists
2. If exists and permissions allow, truncates to zero length
3. If doesn't exist, creates new file on disk
4. Creates entry in Global File Table
5. Allocates lowest unused FD in process FD table
6. Returns FD or -1 on error

3.2 2. Open System Call

open() - Open or Create a File

```
int open(const char *pathname, int flags, mode_t mode);
```

3.2.1 Function Details

- Header:** <fcntl.h> and <unistd.h>
- Purpose:** Open existing file or create new file
- Flexible:** More flexible than `creat()` (can open for read/write)
- Flags:** Control how file is opened (read/write, create, append, etc.)
- Mode:** Required only when `O_CREAT` is used

3.2.2 Common Open Flags

Flag	Value	Description
<code>O_RDONLY</code>	0	Open for reading only
<code>O_WRONLY</code>	1	Open for writing only
<code>O_RDWR</code>	2	Open for reading and writing
<code>O_CREAT</code>	0100	Create file if it doesn't exist
<code>O_EXCL</code>	0200	With <code>O_CREAT</code> , fail if file exists
<code>O_TRUNC</code>	01000	Truncate file to zero length
<code>O_APPEND</code>	02000	Append to file on each write
<code>O_NONBLOCK</code>	04000	Non-blocking I/O
<code>O_SYNC</code>	010000	Synchronous I/O (wait for write completion)

3.2.3 open() Examples

```

1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <errno.h>
5 #include <string.h>
6
7 extern int errno;
8
9 int main() {
10     int fd;
11
12     // Example 1: Open existing file for reading
13     fd = open("existing.txt", O_RDONLY);
14     if (fd == -1) {
15         printf("Error %d: %s\n", errno, strerror(errno));
16     } else {
17         printf("Opened for reading, fd = %d\n", fd);
18         close(fd);
19 }
```

```

20
21 // Example 2: Create new file (or open if exists) for writing
22 fd = open("data.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644);
23 if (fd == -1) {
24     perror("open failed");
25 } else {
26     printf("Created/opened for writing, fd = %d\n", fd);
27     write(fd, "Hello\n", 6);
28     close(fd);
29 }

30
31 // Example 3: Open for append (adds to end of file)
32 fd = open("log.txt", O_WRONLY | O_CREAT | O_APPEND, 0644);
33 if (fd != -1) {
34     printf("Opened for append, fd = %d\n", fd);
35     write(fd, "Log entry\n", 10);
36     close(fd);
37 }

38
39 // Example 4: Exclusive creation (fails if file exists)
40 fd = open("unique.txt", O_WRONLY | O_CREAT | O_EXCL, 0644);
41 if (fd == -1) {
42     if (errno == EEXIST) {
43         printf("File already exists (exclusive mode)\n");
44     }
45 } else {
46     printf("Exclusively created, fd = %d\n", fd);
47     close(fd);
48 }

49
50 return 0;
51 }
```

Listing 4: Various open() Usage Patterns

3.2.4 Path Types

Absolute vs Relative Paths

- **Absolute Path:** Full path from root directory
 - `open("/home/user/file.txt", ...)`
 - Always works from any directory
- **Relative Path:** Path relative to current working directory
 - `open("file.txt", ...)` (current directory)
 - `open("../parent/file.txt", ...)` (parent directory)
 - Depends on where program is executed

3.3 3. Close System Call

close() - Close a File Descriptor

```
int close(int fd);
```

3.3.1 Function Details

- **Header:** <unistd.h>
- **Purpose:** Close file descriptor and release resources
- **Important:** Always close files when done to prevent resource leaks
- **Safe:** close() on already closed FD returns -1 with EBADF
- **Special:** Closing FD 0, 1, or 2 affects standard streams

3.3.2 close() Examples

```
1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5
6 int main() {
7     int fd1, fd2;
8
9     // Open two files
10    fd1 = open("file1.txt", O_RDONLY);
11    if (fd1 == -1) {
12        perror("open file1");
13        return 1;
14    }
15    printf("Opened file1, fd = %d\n", fd1);
16
17    fd2 = open("file2.txt", O_WRONLY | O_CREAT, 0644);
18    if (fd2 == -1) {
19        perror("open file2");
20        close(fd1); // Close first file before returning
21        return 1;
22    }
23    printf("Opened file2, fd = %d\n", fd2);
24
25    // Use the files...
26    char buffer[100];
27    int n = read(fd1, buffer, sizeof(buffer)-1);
28    if (n > 0) {
29        write(fd2, buffer, n);
30    }
31
32    // Close files in reverse order (good practice)
33    if (close(fd2) == -1) {
34        perror("close file2");
35    } else {
36        printf("Closed file2 (fd=%d)\n", fd2);
37    }
38
39    if (close(fd1) == -1) {
```

```

40     perror("close file1");
41 } else {
42     printf("Closed file1 (fd=%d)\n", fd1);
43 }
44
45 // Demonstrate FD reuse
46 int fd3 = open("file3.txt", O_WRONLY | O_CREAT, 0644);
47 printf("Opened file3, fd = %d (reused from closed fds)\n", fd3);
48 close(fd3);
49
50 return 0;
51 }
```

Listing 5: Proper close() Usage

3.3.3 What Happens on close()?

1. Kernel decrements reference count in Global File Table
2. If reference count reaches 0:
 - Flushes any pending writes (if buffered)
 - Updates file metadata (size, modification time)
 - Removes entry from Global File Table
 - For disk files, data is now safely stored
3. Marks FD as available in process FD table
4. Returns 0 on success, -1 on error

3.3.4 Important close() Behavior

```

1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4
5 int main() {
6     // Open same file twice - gets different FDs
7     int fd1 = open("test.txt", O_RDONLY);
8     int fd2 = open("test.txt", O_RDONLY);
9
10    printf("fd1 = %d, fd2 = %d\n", fd1, fd2);
11
12    // Each has independent file position
13    char c;
14    read(fd1, &c, 1);
15    printf("fd1 read: %c\n", c);
16
17    read(fd2, &c, 1);
18    printf("fd2 read: %c (starts from beginning)\n", c);
19
20    // Close fd1 doesn't affect fd2
21    close(fd1);
22    printf("Closed fd1\n");
23
24    // fd2 still works
25    read(fd2, &c, 1);
```

```
26     printf("fd2 still works, read: %c\n", c);
27
28     close(fd2);
29
30     // Try to close invalid FD
31     if (close(999) == -1) {
32         perror("close(999)");
33     }
34
35     // Close stdin (FD 0) - be careful!
36     // close(STDIN_FILENO);
37     // Now read() from stdin will fail!
38
39     return 0;
40 }
```

Listing 6: Understanding close() Behavior

3.4 4. Read System Call

read() - Read from a File Descriptor

```
ssize_t read(int fd, void *buf, size_t count);
```

3.4.1 Function Details

- **Header:** <unistd.h>
- **Purpose:** Read data from file descriptor into buffer
- **Returns:** Number of bytes read, 0 on EOF, -1 on error
- **Blocking:** By default, blocks if no data available
- **Position:** Updates file position (except for pipes, sockets)

3.4.2 read() Examples

```
1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5 #include <errno.h>
6
7 int main() {
8     int fd;
9     char buffer[1024];
10    ssize_t bytes_read;
11
12    // Open file for reading
13    fd = open("input.txt", O_RDONLY);
14    if (fd == -1) {
15        perror("open");
16        return 1;
17    }
18
19    printf("Reading file (fd=%d)... \n", fd);
20
21    // Read file in chunks
22    int total_bytes = 0;
23    while ((bytes_read = read(fd, buffer, sizeof(buffer))) > 0) {
24        total_bytes += bytes_read;
25        printf("Read %zd bytes (total: %d)\n", bytes_read, total_bytes);
26
27        // Process the buffer (e.g., write to stdout)
28        write(STDOUT_FILENO, buffer, bytes_read);
29    }
30
31    // Check why loop ended
32    if (bytes_read == 0) {
33        printf("Reached end of file (EOF)\n");
34    } else if (bytes_read == -1) {
35        perror("read error");
36    }
37
38    close(fd);
```

```

40 // Example 2: Reading from stdin
41 printf("\nEnter something (press Ctrl+D for EOF):\n");
42 bytes_read = read(STDIN_FILENO, buffer, sizeof(buffer)-1);
43
44 if (bytes_read > 0) {
45     buffer[bytes_read] = '\0';
46     printf("You entered: %s", buffer);
47 } else if (bytes_read == 0) {
48     printf("EOF on stdin\n");
49 }
50
51 return 0;
52 }
```

Listing 7: Using read() System Call

3.4.3 Important read() Behaviors

```

1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5
6 int main() {
7     // Create a test file
8     int fd = open("test.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644);
9     write(fd, "ABCDEFGHIJ", 10);
10    close(fd);
11
12    // Open two file descriptors to same file
13    int fd1 = open("test.txt", O_RDONLY);
14    int fd2 = open("test.txt", O_RDONLY);
15
16    char c;
17
18    // Each FD has independent file position
19    read(fd1, &c, 1);
20    printf("fd1 read: %c\n", c); // A
21
22    read(fd2, &c, 1);
23    printf("fd2 read: %c\n", c); // A (independent position)
24
25    read(fd1, &c, 1);
26    printf("fd1 read: %c\n", c); // B
27
28    // Using lseek to change position
29    lseek(fd1, 4, SEEK_SET); // Move to position 4 (0-based)
30    read(fd1, &c, 1);
31    printf("fd1 after lseek: %c\n", c); // E
32
33    // fd2 position unchanged
34    read(fd2, &c, 1);
35    printf("fd2 read: %c\n", c); // B
36
37    close(fd1);
38    close(fd2);
39
40    // Example: Partial read
41    fd = open("test.txt", O_RDONLY);
42    char buf[5];
```

```

43     ssize_t n;
44
45     // Request 5 bytes, might get less
46     n = read(fd, buf, 3); // Only ask for 3
47     printf("Requested 3, got %zd bytes\n", n);
48
49     // Request more than available
50     n = read(fd, buf, 20); // File has 7 bytes left
51     printf("Requested 20, got %zd bytes (until EOF)\n", n);
52
53     // At EOF
54     n = read(fd, buf, 5);
55     printf("At EOF: requested 5, got %zd bytes\n", n); // 0
56
57     close(fd);
58
59     return 0;
60 }
```

Listing 8: Understanding read() Behavior

3.4.4 read() Return Value Summary

Return Value	Meaning
<code>n > 0</code>	Successfully read <code>n</code> bytes
<code>0</code>	End of File (EOF) reached
<code>-1</code>	Error occurred (check <code>errno</code>)
<code>0 < n < count</code>	Partial read (common with pipes, sockets)

3.4.5 Common read() Errors

- `EBADF`: `fd` is not a valid file descriptor
- `EINTR`: Interrupted by signal before reading any data
- `EINVAL`: `fd` is attached to object unsuitable for reading
- `EIO`: I/O error (low-level hardware error)
- `EISDIR`: `fd` refers to a directory

3.5 5. Write System Call

write() - Write to a File Descriptor

```
ssize_t write(int fd, const void *buf, size_t count);
```

3.5.1 Function Details

- **Header:** <unistd.h>
- **Purpose:** Write data from buffer to file descriptor
- **Returns:** Number of bytes written, -1 on error
- **Blocking:** May block if device/buffer is full
- **Atomic:** Write up to PIPE_BUF bytes (usually 4096) is atomic

3.5.2 write() Examples

```
1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5 #include <string.h>
6
7 int main() {
8     int fd;
9     ssize_t bytes_written;
10    const char *text = "Hello, System Calls!\n";
11
12    // Example 1: Write to a file
13    fd = open("output.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644);
14    if (fd == -1) {
15        perror("open");
16        return 1;
17    }
18
19    bytes_written = write(fd, text, strlen(text));
20    printf("Wrote %zd bytes to file\n", bytes_written);
21
22    // Write more data
23    const char *more_text = "Second line of text.\n";
24    bytes_written = write(fd, more_text, strlen(more_text));
25    printf("Wrote %zd more bytes\n", bytes_written);
26
27    close(fd);
28
29    // Example 2: Append to file
30    fd = open("output.txt", O_WRONLY | O_APPEND);
31    const char *append_text = "Appended line.\n";
32    write(fd, append_text, strlen	append_text));
33    close(fd);
34
35    // Example 3: Write to stdout and stderr
36    write(STDOUT_FILENO, "This goes to stdout\n", 20);
37    write(STDERR_FILENO, "This goes to stderr\n", 20);
38
39    // Example 4: Copy file using read/write
```

```

40     int src_fd = open("input.txt", O_RDONLY);
41     int dst_fd = open("copy.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644);
42
43     if (src_fd != -1 && dst_fd != -1) {
44         char buffer[4096];
45         ssize_t n;
46
47         printf("Copying file...\n");
48         while ((n = read(src_fd, buffer, sizeof(buffer))) > 0) {
49             ssize_t written = write(dst_fd, buffer, n);
50             if (written != n) {
51                 perror("write error during copy");
52                 break;
53             }
54         }
55
56         if (n == -1) {
57             perror("read error during copy");
58         } else {
59             printf("File copied successfully\n");
60         }
61
62         close(src_fd);
63         close(dst_fd);
64     }
65
66     return 0;
67 }
```

Listing 9: Using write() System Call

3.5.3 Important write() Behaviors

```

1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5 #include <string.h>
6
7 int main() {
8     // Example 1: Partial writes
9     int fd = open("partial.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644);
10    char large_buffer[10000];
11    memset(large_buffer, 'A', sizeof(large_buffer));
12
13    // write() may write less than requested
14    ssize_t written = write(fd, large_buffer, sizeof(large_buffer));
15    printf("Requested %zu, wrote %zd bytes\n", sizeof(large_buffer),
16          written);
17    close(fd);
18
19    // Example 2: O_APPEND behavior
20    fd = open("append.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644);
21    write(fd, "Line 1\n", 7);
22    close(fd);
23
24    // Open with O_APPEND - all writes go to end
25    int fd1 = open("append.txt", O_WRONLY | O_APPEND);
26    int fd2 = open("append.txt", O_WRONLY); // Without O_APPEND
```

```

27 // Both write at current position
28 write(fd1, "fd1: Line 2\n", 12); // Goes to end (due to O_APPEND)
29
30 lseek(fd2, 0, SEEK_SET); // Move to beginning
31 write(fd2, "OVERWRITE", 9); // Overwrites beginning!
32
33 close(fd1);
34 close(fd2);
35
36 // Example 3: Writing to pipes
37 int pipefd[2];
38 pipe(pipefd);
39
40 char msg[] = "Hello through pipe!";
41 written = write(pipefd[1], msg, sizeof(msg));
42 printf("Wrote %zd bytes to pipe\n", written);
43
44 // Pipe has limited buffer (usually 64KB)
45 // Write may block if pipe is full
46
47 close(pipefd[0]);
48 close(pipefd[1]);
49
50 // Example 4: Error cases
51 fd = open("/dev/full", O_WRONLY); // Device that always returns ENOSPC
52 if (fd != -1) {
53     written = write(fd, "test", 4);
54     if (written == -1) {
55         perror("write to /dev/full");
56     }
57     close(fd);
58 }
59
60 // Write to closed FD
61 fd = 999; // Invalid FD
62 if (write(fd, "test", 4) == -1) {
63     perror("write to invalid fd");
64 }
65
66 return 0;
67 }
```

Listing 10: Understanding write() Behavior

3.5.4 write() Return Value Summary

Return Value	Meaning
n == count	All bytes successfully written
0 < n < count	Partial write (may retry remaining)
-1	Error occurred (check errno)
0	Rare (usually means count was 0)

3.5.5 Common write() Errors

- EBADF: fd is not a valid file descriptor open for writing
- EFBIG: File size limit exceeded

- **EINTR:** Interrupted by signal before writing any data
- **EIO:** Low-level I/O error
- **ENOSPC:** No space left on device
- **EPIPE:** Broken pipe (reader closed)

4 Practical Examples and Use Cases

4.1 Example 1: Simple File Copy

```
1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5
6 #define BUFFER_SIZE 4096
7
8 int main(int argc, char *argv[]) {
9     if (argc != 3) {
10         fprintf(stderr, "Usage: %s <source> <destination>\n", argv[0]);
11         return 1;
12     }
13
14     const char *source = argv[1];
15     const char *dest = argv[2];
16
17     // Open source file
18     int src_fd = open(source, O_RDONLY);
19     if (src_fd == -1) {
20         perror("Error opening source file");
21         return 1;
22     }
23
24     // Open destination file
25     int dest_fd = open(dest, O_WRONLY | O_CREAT | O_TRUNC, 0644);
26     if (dest_fd == -1) {
27         perror("Error opening destination file");
28         close(src_fd);
29         return 1;
30     }
31
32     // Copy data
33     char buffer[BUFFER_SIZE];
34     ssize_t bytes_read, bytes_written;
35     int total_bytes = 0;
36
37     while ((bytes_read = read(src_fd, buffer, BUFFER_SIZE)) > 0) {
38         bytes_written = write(dest_fd, buffer, bytes_read);
39
40         if (bytes_written != bytes_read) {
41             perror("Error writing to destination");
42             close(src_fd);
43             close(dest_fd);
44             return 1;
45         }
46
47         total_bytes += bytes_written;
48     }
49
50     if (bytes_read == -1) {
51         perror("Error reading from source");
52     } else {
53         printf("Copied %d bytes from '%s' to '%s'\n",
54               total_bytes, source, dest);
55     }
56 }
```

```

57     // Cleanup
58     close(src_fd);
59     close(dest_fd);
60
61     return 0;
62 }
```

Listing 11: File Copy Utility Using System Calls

4.2 Example 2: Simple Cat Command

```

1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5
6 #define BUFFER_SIZE 1024
7
8 void cat_file(const char *filename) {
9     int fd;
10
11     if (filename[0] == '-' && filename[1] == '\0') {
12         // Read from stdin
13         fd = STDIN_FILENO;
14     } else {
15         fd = open(filename, O_RDONLY);
16         if (fd == -1) {
17             perror(filename);
18             return;
19         }
20     }
21
22     char buffer[BUFFER_SIZE];
23     ssize_t bytes_read;
24
25     while ((bytes_read = read(fd, buffer, BUFFER_SIZE)) > 0) {
26         write(STDOUT_FILENO, buffer, bytes_read);
27     }
28
29     if (bytes_read == -1) {
30         perror("read error");
31     }
32
33     if (fd != STDIN_FILENO) {
34         close(fd);
35     }
36 }
37
38 int main(int argc, char *argv[]) {
39     if (argc == 1) {
40         // No arguments, read from stdin
41         cat_file("-");
42     } else {
43         for (int i = 1; i < argc; i++) {
44             cat_file(argv[i]);
45         }
46     }
47
48     return 0;
49 }
```

Listing 12: Simple cat command implementation

4.3 Example 3: Simple Echo Command

```

1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <string.h>
5 #include <stdlib.h>
6
7 int main(int argc, char *argv[]) {
8     int output_fd = STDOUT_FILENO;
9     int append_mode = 0;
10
11    // Check for output redirection
12    for (int i = 1; i < argc; i++) {
13        if (strcmp(argv[i], ">") == 0 && i + 1 < argc) {
14            // Truncate and write
15            output_fd = open(argv[i + 1],
16                              O_WRONLY | O_CREAT | O_TRUNC, 0644);
17            if (output_fd == -1) {
18                perror("open for output");
19                return 1;
20            }
21            // Remove redirection tokens from args
22            argc = i;
23            break;
24        } else if (strcmp(argv[i], ">>") == 0 && i + 1 < argc) {
25            // Append mode
26            output_fd = open(argv[i + 1],
27                              O_WRONLY | O_CREAT | O_APPEND, 0644);
28            if (output_fd == -1) {
29                perror("open for append");
30                return 1;
31            }
32            argc = i;
33            append_mode = 1;
34            break;
35        }
36    }
37
38    // Write all arguments
39    for (int i = 1; i < argc; i++) {
40        write(output_fd, argv[i], strlen(argv[i]));
41        if (i < argc - 1) {
42            write(output_fd, " ", 1);
43        }
44    }
45
46    write(output_fd, "\n", 1);
47
48    if (output_fd != STDOUT_FILENO) {
49        close(output_fd);
50        if (append_mode) {
51            printf("Appended to file\n");
52        } else {
53            printf("Output written to file\n");
54        }
55    }
56}
```

```

55     }
56
57     return 0;
58 }
```

Listing 13: Echo command with redirection support

4.4 Example 4: Simple Head Command

```

1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5
6 #define DEFAULT_LINES 10
7
8 void head_file(const char *filename, int line_count) {
9     int fd;
10
11     if (filename[0] == '-' && filename[1] == '\0') {
12         fd = STDIN_FILENO;
13     } else {
14         fd = open(filename, O_RDONLY);
15         if (fd == -1) {
16             perror(filename);
17             return;
18         }
19     }
20
21     char buffer[1];
22     int lines_printed = 0;
23     ssize_t bytes_read;
24     char last_char = '\n'; // Start as if we just saw a newline
25
26     while (lines_printed < line_count &&
27            (bytes_read = read(fd, buffer, 1)) > 0) {
28
29         write(STDOUT_FILENO, buffer, 1);
30
31         if (buffer[0] == '\n') {
32             lines_printed++;
33         }
34
35         last_char = buffer[0];
36     }
37
38     if (last_char != '\n') {
39         write(STDOUT_FILENO, "\n", 1);
40     }
41
42     if (fd != STDIN_FILENO) {
43         close(fd);
44     }
45 }
46
47 int main(int argc, char *argv[]) {
48     int line_count = DEFAULT_LINES;
49     int file_index = 1;
50
51     // Parse command line arguments
```

```
52     if (argc > 1 && argv[1][0] == '-') {
53         line_count = atoi(&argv[1][1]);
54         if (line_count <= 0) {
55             line_count = DEFAULT_LINES;
56         }
57         file_index = 2;
58     }
59
60     if (file_index >= argc) {
61         // No filename, read from stdin
62         head_file("-", line_count);
63     } else {
64         for (int i = file_index; i < argc; i++) {
65             if (argc - file_index > 1) {
66                 // Multiple files, print header
67                 printf("\n==> %s <==\n", argv[i]);
68             }
69             head_file(argv[i], line_count);
70         }
71     }
72
73     return 0;
74 }
```

Listing 14: Head command showing first N lines

5 Advanced Topics and Best Practices

5.1 Error Handling with `errno`

Using `errno` for Error Diagnostics

The `errno` variable is set by system calls on error. Always check it after a failed call.

```
1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5 #include <errno.h>
6 #include <string.h>
7
8 int main() {
9     int fd;
10
11     // Example 1: Simple error check
12     fd = open("nonexistent.txt", O_RDONLY);
13     if (fd == -1) {
14         printf("Error %d: %s\n", errno, strerror(errno));
15     }
16
17     // Example 2: perror() for automatic message
18     fd = open("/root/protected.txt", O_RDONLY);
19     if (fd == -1) {
20         perror("open failed");
21         // Prints: open failed: Permission denied
22     }
23
24     // Example 3: Checking specific errors
25     fd = open("test.txt", O_RDONLY | O_CREAT | O_EXCL, 0644);
26     if (fd == -1) {
27         if (errno == EEXIST) {
28             printf("File already exists (exclusive mode)\n");
29         } else if (errno == EACCES) {
30             printf("Permission denied\n");
31         } else {
32             perror("unexpected error");
33         }
34     }
35
36     // Example 4: Save/Restore errno for library calls
37     int saved_errno;
38     fd = open("file.txt", O_RDONLY);
39     if (fd == -1) {
40         saved_errno = errno;
41         // Call some library function that might change errno
42         printf("Some message\n");
43         // Restore errno for accurate error reporting
44         errno = saved_errno;
45         perror("open");
46     }
47
48     return 0;
49 }
```

Listing 15: Proper Error Handling

5.2 File Positioning with lseek

lseek() - Reposition Read/Write Offset

```
off_t lseek(int fd, off_t offset, int whence);
```

5.2.1 lseek() Examples

```

1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5
6 int main() {
7     int fd = open("test.txt", O_RDWR | O_CREAT | O_TRUNC, 0644);
8     if (fd == -1) {
9         perror("open");
10        return 1;
11    }
12
13    // Write some data
14    write(fd, "0123456789ABCDEF", 16);
15
16    // Move to different positions
17    off_t pos;
18    char buffer[2];
19
20    // SEEK_SET: offset from beginning
21    lseek(fd, 5, SEEK_SET);
22    read(fd, buffer, 1);
23    buffer[1] = '\0';
24    printf("Byte at position 5: %s\n", buffer); // 5
25
26    // SEEK_CUR: offset from current position
27    lseek(fd, 3, SEEK_CUR); // Move 3 bytes forward
28    read(fd, buffer, 1);
29    printf("3 bytes after position 5: %s\n", buffer); // 9
30
31    // SEEK_END: offset from end
32    lseek(fd, -4, SEEK_END); // 4 bytes from end
33    read(fd, buffer, 1);
34    printf("4 bytes from end: %s\n", buffer); // C
35
36    // Get current position
37    pos = lseek(fd, 0, SEEK_CUR);
38    printf("Current position: %ld\n", pos);
39
40    // Get file size
41    off_t size = lseek(fd, 0, SEEK_END);
42    printf("File size: %ld bytes\n", size);
43
44    // Rewind to beginning
45    lseek(fd, 0, SEEK_SET);
46
```

```

47 // Write at specific position (overwrites)
48 lseek(fd, 3, SEEK_SET);
49 write(fd, "XXX", 3);
50
51 // Read modified file
52 lseek(fd, 0, SEEK_SET);
53 char all[20];
54 int n = read(fd, all, 16);
55 all[n] = '\0';
56 printf("Modified content: %s\n", all); // 012XXX6789ABCDEF
57
58 close(fd);
59 return 0;
60 }
```

Listing 16: Using lseek() for File Positioning

5.3 File Status with fstat

fstat() - Get File Status

```
int fstat(int fd, struct stat *statbuf);
```

```

1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5 #include <sys/stat.h>
6 #include <time.h>
7
8 int main() {
9     int fd = open("test.txt", O_RDONLY);
10    if (fd == -1) {
11        perror("open");
12        return 1;
13    }
14
15    struct stat file_stat;
16
17    if (fstat(fd, &file_stat) == -1) {
18        perror("fstat");
19        close(fd);
20        return 1;
21    }
22
23    printf("File Information:\n");
24    printf("  File size: %ld bytes\n", file_stat.st_size);
25    printf("  Blocks allocated: %ld\n", file_stat.st_blocks);
26    printf("  Block size: %ld bytes\n", file_stat.st_blksize);
27    printf("  Inode number: %ld\n", file_stat.st_ino);
28    printf("  Links: %ld\n", file_stat.st_nlink);
29    printf("  Owner UID: %d\n", file_stat.st_uid);
30    printf("  Group GID: %d\n", file_stat.st_gid);
31
32    printf("  Permissions: ");
33    printf((S_ISDIR(file_stat.st_mode)) ? "d" : "-");
34    printf((file_stat.st_mode & S_IRUSR) ? "r" : "-");
35    printf((file_stat.st_mode & S_IWUSR) ? "w" : "-");
36    printf((file_stat.st_mode & S_IXUSR) ? "x" : "-");
```

```
37     printf((file_stat.st_mode & S_IRGRP) ? "r" : "-");
38     printf((file_stat.st_mode & S_IWGRP) ? "w" : "-");
39     printf((file_stat.st_mode & S_IXGRP) ? "x" : "-");
40     printf((file_stat.st_mode & S_IROTH) ? "r" : "-");
41     printf((file_stat.st_mode & S_IWOTH) ? "w" : "-");
42     printf((file_stat.st_mode & S_IXOTH) ? "x" : "-");
43     printf("\n");
44
45     printf("  Last access: %s", ctime(&file_stat.st_atime));
46     printf("  Last modification: %s", ctime(&file_stat.st_mtime));
47     printf("  Last status change: %s", ctime(&file_stat.st_ctime));
48
49     close(fd);
50     return 0;
51 }
```

Listing 17: Getting File Information with fstat()

5.4 Duplicating File Descriptors with dup/dup2

dup()/dup2() - Duplicate File Descriptors

```
int dup(int oldfd);
int dup2(int oldfd, int newfd);
```

```

1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5
6 int main() {
7     int fd, fd2, fd3;
8
9     // Create a file
10    fd = open("test.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644);
11    write(fd, "Original FD\n", 12);
12
13    // dup() creates a copy with lowest available FD
14    fd2 = dup(fd);
15    printf("Original fd: %d, dup fd: %d\n", fd, fd2);
16
17    // Both write to same file
18    write(fd, "From original fd\n", 17);
19    write(fd2, "From dup fd\n", 12);
20
21    // They share file position
22    lseek(fd, 0, SEEK_SET);
23    char buffer[100];
24    int n = read(fd, buffer, sizeof(buffer)-1);
25    buffer[n] = '\0';
26    printf("File content:\n%s", buffer);
27
28    // dup2() duplicates to specific FD
29    fd3 = dup2(fd, 10); // Try to duplicate to FD 10
30    printf("dup2 created fd: %d\n", fd3);
31
32    // Redirect stdout to file
33    int saved_stdout = dup(STDOUT_FILENO); // Save original stdout
34    dup2(fd, STDOUT_FILENO); // Now stdout goes to file
35
36    printf("This goes to the file, not screen!\n");
37
38    // Restore stdout
39    dup2(saved_stdout, STDOUT_FILENO);
40    printf("Back to screen output\n");
41
42    // Cleanup
43    close(fd);
44    close(fd2);
45    close(fd3);
46    close(saved_stdout);
47
48    return 0;
49 }
```

Listing 18: Duplicating File Descriptors

5.5 Non-blocking I/O

```

1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5 #include <errno.h>
6
7 int main() {
8     // Set stdin to non-blocking mode
9     int flags = fcntl(STDIN_FILENO, F_GETFL, 0);
10    fcntl(STDIN_FILENO, F_SETFL, flags | O_NONBLOCK);
11
12    printf("Type something (non-blocking read, Ctrl+D to exit)...\\n");
13
14    char buffer[100];
15    int bytes_read;
16
17    for (int i = 0; i < 10; i++) {
18        bytes_read = read(STDIN_FILENO, buffer, sizeof(buffer)-1);
19
20        if (bytes_read > 0) {
21            buffer[bytes_read] = '\\0';
22            printf("You typed: %s", buffer);
23            break;
24        } else if (bytes_read == -1 && errno == EAGAIN) {
25            // No data available yet
26            printf(".");
27            fflush(stdout);
28            sleep(1);
29        } else if (bytes_read == 0) {
30            printf("\\nEOF reached\\n");
31            break;
32        }
33    }
34
35    if (bytes_read <= 0 && errno != EAGAIN) {
36        printf("\\nNo input received\\n");
37    }
38
39    // Restore blocking mode
40    fcntl(STDIN_FILENO, F_SETFL, flags);
41
42    return 0;
43 }
```

Listing 19: Non-blocking I/O Example

6 Common Pitfalls and Debugging

6.1 Common System Call Errors

Error	Cause	Solution
Memory Leaks	Not closing file descriptors	Always close() when done
Race Conditions	File created between check and open	Use O_CREAT O_EXCL
Permission Issues	Wrong file permissions	Check umask, use correct mode
Buffer Overflows	Not checking read() return value	Always check return values
Partial Reads/Writes	Assuming full buffer transferred	Handle partial operations
File Descriptor Limits	Too many open files	Close unused FDs, increase limits
Blocking Forever	Read/write on blocking FD	Use non-blocking or timeouts

6.2 Debugging Techniques

```

1 #include <fcntl.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5 #include <errno.h>
6
7 // Debug macro
8 #ifdef DEBUG
9 #define DBG_PRINT(fmt, ...) fprintf(stderr, "DEBUG: " fmt, ##__VA_ARGS__)
10#else
11#define DBG_PRINT(fmt, ...)
12#endif
13
14 int safe_open(const char *path, int flags, mode_t mode) {
15    DBG_PRINT("Opening %s with flags %d\n", path, flags);
16
17    int fd = open(path, flags, mode);
18    if (fd == -1) {
19        DBG_PRINT("open failed: %s\n", strerror(errno));
20    } else {
21        DBG_PRINT("open succeeded, fd = %d\n", fd);
22    }
23
24    return fd;
25}
26
27 ssize_t safe_read(int fd, void *buf, size_t count) {
28    DBG_PRINT("Reading from fd %d, requesting %zu bytes\n", fd, count);
29
30    ssize_t result = read(fd, buf, count);
31    if (result == -1) {
32        DBG_PRINT("read failed: %s\n", strerror(errno));

```

```
33     } else if (result == 0) {
34         DBG_PRINT("read reached EOF\n");
35     } else {
36         DBG_PRINT("read returned %zd bytes\n", result);
37     }
38
39     return result;
40 }
41
42 int main() {
43     // Enable debugging
44     #define DEBUG 1
45
46     int fd = safe_open("test.txt", O_RDONLY, 0);
47     if (fd == -1) {
48         return 1;
49     }
50
51     char buffer[100];
52     ssize_t n = safe_read(fd, buffer, sizeof(buffer));
53
54     if (n > 0) {
55         write(STDOUT_FILENO, buffer, n);
56     }
57
58     close(fd);
59
60     // Check file descriptor limits
61     #ifdef DEBUG
62     printf("\nFile descriptor information:\n");
63     printf("STDIN_FILENO = %d\n", STDIN_FILENO);
64     printf("STDOUT_FILENO = %d\n", STDOUT_FILENO);
65     printf("STDERR_FILENO = %d\n", STDERR_FILENO);
66
67     // Try to find maximum FDs
68     for (int i = 3; i < 100; i++) {
69         if (fcntl(i, F_GETFD) == -1 && errno == EBADF) {
70             printf("First unused FD: %d\n", i);
71             break;
72         }
73     }
74     #endif
75
76     return 0;
77 }
```

Listing 20: Debugging System Call Programs

6.3 Best Practices Summary

System Calls Best Practices

- 1. Always check return values** - Never assume success
- 2. Use proper error handling** - Check errno, use perror()
- 3. Close file descriptors** - Prevent resource leaks
- 4. Handle partial reads/writes** - Loop until all data transferred
- 5. Use appropriate flags** - O_CREAT, O_EXCL for atomic operations
- 6. Check permissions** - Use correct mode values
- 7. Be aware of blocking** - Consider non-blocking I/O for responsiveness
- 8. Use constants** - STDIN_FILENO instead of 0, etc.
- 9. Test edge cases** - Empty files, large files, permission issues
- 10. Clean up on failure** - Close opened FDs before returning

6.4 Performance Considerations

- Buffer size matters:** Too small = many system calls, too large = memory waste
- System calls are expensive:** Minimize context switches
- Use read/write efficiently:** Read/write in reasonable chunks (4KB-64KB)
- Consider memory mapping:** mmap() for large files
- Batch operations:** Combine when possible
- Avoid unnecessary seeks:** Sequential access is faster

Conclusion

I/O system calls are the foundation of file operations in Linux/Unix systems. Understanding them provides deep insight into how operating systems work and enables writing efficient, low-level code.

Key Takeaways

- **File descriptors** are integers that represent open files in a process
- **Standard streams** (0,1,2) are automatically opened for every process
- The **5 essential I/O system calls** are: create, open, close, read, write
- **Always check return values** and handle errors properly
- **System calls provide low-level control** but require careful management
- **Understanding these concepts** is crucial for system programming

Happy Systems Programming!

Master the fundamentals, control the system!

This document was compiled on January 15, 2026.

For updates and more resources, visit: <https://github.com/systems-programming>